

Sukriti Srivastava

(240)-714-8344 | srivastavasukriti2000@gmail.com | [LinkedIn](#) | [GitHub](#)

EDUCATION

University of Maryland

Master of Science in Data Science | GPA: 4.0

Coursework: Big Data Systems, Algorithms for Data Science, GenAI, Agentic AI, Principles of Machine Learning

College Park, MD

August 2025 – May 2027

Manipal Institute of Technology

Bachelor of Technology in Electronics and Communication Engineering | CGPA: 8.68

Manipal, India

July 2018 – July 2022

WORK EXPERIENCE

Software Engineer (Data Engineer)

July 2022 – August 2025

NatWest Group

Bangalore, India

- Built a centralized data pipeline framework on AWS SageMaker Unified Studio with **8** reusable templates and CloudWatch monitoring, reducing pipeline setup time by **70%**.
- Developed a metadata-driven Spark/Scala framework for batch and streaming workflows with Kafka-triggered Airflow orchestration, enabling teams to build new pipelines through configuration instead of custom code.
- Integrated OpenLineage into the Spark/Scala framework, enabling end-to-end data lineage tracking across all production pipelines and reducing pipeline debugging time by **80%**.
- Automated pipeline deployments across **5 environments** using GitLab CI/CD and PostgreSQL stored procedures for environment promotion, reducing manual configuration effort by 90%.
- Led the migration of **200+** ETL pipelines onto the Spark/Scala framework, partnering with cross-functional teams to standardize execution and authoring configuration documentation that reduced pipeline onboarding time by **60%**.
- Engineered **50+** Informatica-based ETL pipelines to migrate data from AWS S3 to Snowflake, enforcing validation across **6** data quality dimensions to deliver clean, analytics-ready datasets.

PROJECTS

F1-Analytics: Formula 1 Lakehouse | Databricks, PySpark, Delta Lake, SQL

- Built an end-to-end lakehouse on Databricks using medallion architecture with incremental loading, ingesting **77 seasons** of Formula 1 race data with Unity Catalog governance and Lakeflow Jobs orchestration.
- Designed gold-layer data models and **6 analytical views** covering driver and constructor standings and teammate head-to-head stats, powering an AI/BI dashboard.
- Developed a scoring model to predict 2026 standings from 5 seasons of results, using exponential decay to prioritize recent form and a median baseline fallback for new teams.

HealthShield: Health Insurance Cost & Fraud Analytics | Python, SQL, SQLite, scikit-learn, Streamlit, Git

- Built a reproducible three-layer data pipeline over **558K Medicare claims (5,410 providers)** and **100K member records**, with automated leakage checks that caught a target-leaking feature before it reached modeling.
- Developed a provider fraud model where **84% of the top 5%** of flagged providers were fraudulent vs. 9.4% at random, outperforming a billing-volume baseline.
- Trained a member cost model for budgeting that predicted group-level total costs within **0.79%** of actual, with results displayed in a Streamlit dashboard.

ClariCase: Consumer Complaint Routing | Python, scikit-learn, Supabase (PostgreSQL), Streamlit

- Built a deployed platform that categorizes free-text consumer complaints by issue and routes them to **11 support teams**, trained on a **1.9 GB** federal complaint dataset with time-based splits and Supabase-backed tracking.
- Auto-routed **46.5%** of complaints at **94.4% precision** on a holdout of 109,834 complaints.

TECHNICAL SKILLS

Languages	: Python, Scala, SQL, Java, JavaScript
Databases & Warehouses	: PostgreSQL, MySQL, Snowflake, DynamoDB
Big Data & Streaming	: Apache Spark, PySpark, Kafka, Delta Lake
Data Platforms & Tools	: Databricks, Airflow, Informatica
Cloud & DevOps	: AWS (S3, Athena, EC2, CloudWatch, SageMaker), GitLab CI/CD, Docker, Git, Linux
Visualization	: Streamlit, Tableau, Databricks AI/BI Dashboards

CERTIFICATIONS

Databricks Certified Data Engineer Associate

October 2026